

Kevin Hance
DBMS (CPSC 321)
HW7

READING ASSIGNMENT:

1. The main difference between SQL WHERE and HAVING is that WHERE operates on each individual row, whereas HAVING filters based on groups after a GROUP BY statement has been made. This impacts performance, as using a WHERE statement with a GROUP BY will attempt to filter all rows, some of which would not otherwise be included in the query result.
2. All aggregate functions except COUNT(*) will ignore null values. COUNT(*) is the exception in this case because, if there is no CASE statement that the row satisfies the criteria of, the row will be treated as a null value, despite the fact that it may have some non-null values that are worth counting.
3. DROP TABLE IF EXISTS border;
DROP TABLE IF EXISTS city;
DROP TABLE IF EXISTS province;
DROP TABLE IF EXISTS country;

```
CREATE TABLE country(  
    country_code VARCHAR(10),  
    country_name VARCHAR(50),  
    gdp INT UNSIGNED,  
    inflation FLOAT,  
    PRIMARY KEY (country_code)  
);
```

```
CREATE TABLE province(  
    province_name VARCHAR(50),  
    country_code VARCHAR(10),  
    area FLOAT UNSIGNED,  
    PRIMARY KEY (province_name, country_code),  
    FOREIGN KEY (country_code) REFERENCES country (country_code)  
);
```

```
CREATE TABLE city(  
    city_name VARCHAR(50),  
    province_name VARCHAR(50),  
    country_code VARCHAR(10),  
    population INT UNSIGNED,  
    PRIMARY KEY (city_name, province_name, country_code),
```

```

        FOREIGN KEY (province_name, country_code) REFERENCES province
        (province_name, country_code)
    );

```

```

CREATE TABLE border(
    country_code_1 VARCHAR(10),
    country_code_2 VARCHAR(10),
    border_length FLOAT UNSIGNED,
    PRIMARY KEY (country_code_1, country_code_2),
    FOREIGN KEY (country_code_1) REFERENCES country (country_code),
    FOREIGN KEY (country_code_2) REFERENCES country (country_code)
);

```

```

INSERT INTO country VALUES ('OS', 'Oswaldo', 78000, 6.7);
INSERT INTO country VALUES ('RL', 'Renlandia', 54000, 7.7);
INSERT INTO country VALUES ('GN', 'Geneva', 65000, null);

```

```

SELECT count(c.inflation)
FROM country c
GROUP BY country_code
HAVING count(DISTINCT c.country_code) > 0.0;

```

```

mysql>
mysql> SELECT count(c.inflation)
-> FROM country c
-> GROUP BY country_code
-> HAVING count(DISTINCT c.country_code) > 0.0;
+-----+
| count(c.inflation) |
+-----+
| 0 |
| 1 |
| 1 |
+-----+
3 rows in set (0.00 sec)

```

4. A view is similar to a table in that you can query the view for information in the form of tables. A prepared statement is a pre-written query on a table or view that returns a specific result.

TECHNICAL WORK:

1.

```
mysql> -- question 1
mysql> SELECT c.country_name, c.country_code, c.GDP, c.inflation, sum(cit.population)
-> FROM country c JOIN city cit ON cit.country_code = c.country_code
-> GROUP BY c.country_code;

+-----+-----+-----+-----+-----+
| country_name | country_code | GDP   | inflation | sum(cit.population) |
+-----+-----+-----+-----+-----+
| Geneva      | GN          | 65000 | 2.1       | 10834320             |
| Oswaldo     | OS          | 78000 | 6.7       | 728142               |
| Renlandia   | RL          | 54000 | 7.7       | 202026               |
+-----+-----+-----+-----+-----+
3 rows in set (0.01 sec)
```

```
SELECT c.country_name, c.country_code, c.GDP, c.inflation,
sum(cit.population)
FROM country c JOIN city cit ON cit.country_code = c.country_code
GROUP BY c.country_code;
```

2.

```
mysql> SELECT p.country_code, p.province_name, p.area, sum(c.population)
-> FROM province p JOIN city c ON p.province_name = c.province_name AND p.country_code = c.country_code
-> GROUP BY p.province_name
-> HAVING sum(c.population) > @population;

+-----+-----+-----+-----+
| country_code | province_name | area  | sum(c.population) |
+-----+-----+-----+-----+
| GN          | Antalens     | 72003 | 5400120           |
| GN          | Huport       | 54041 | 5530210           |
+-----+-----+-----+-----+
2 rows in set (0.00 sec)
```

```
set @population = 5000000;
SELECT p.country_code, p.province_name, p.area, sum(c.population)
FROM province p JOIN city c ON p.province_name = c.province_name
AND p.country_code = c.country_code
GROUP BY p.province_name
HAVING sum(c.population) > @population;
```

3.

```
mysql> SELECT c.country_code, c.country_name, count(DISTINCT cit.city_name)
-> FROM country c JOIN city cit ON c.country_code = cit.country_code
-> GROUP BY c.country_code
-> ORDER BY count(DISTINCT cit.city_name) DESC;
+-----+-----+-----+
| country_code | country_name | count(DISTINCT cit.city_name) |
+-----+-----+-----+
| RL           | Renlandia    | 5                               |
| OS           | Oswaldo      | 5                               |
| GN           | Geneva       | 4                               |
+-----+-----+-----+
3 rows in set (0.00 sec)
```

```
SELECT c.country_code, c.country_name, count(DISTINCT cit.city_name)
FROM country c JOIN city cit ON c.country_code = cit.country_code
GROUP BY c.country_code
ORDER BY count(DISTINCT cit.city_name) DESC;
```

4.

```
mysql> SELECT c.country_code, c.country_name, sum(p.area)
-> FROM country c JOIN province p ON c.country_code = p.country_code
-> GROUP BY c.country_code
-> ORDER BY sum(p.area) DESC;
+-----+-----+-----+
| country_code | country_name | sum(p.area) |
+-----+-----+-----+
| GN           | Geneva       | 126044       |
| RL           | Renlandia    | 96720        |
| OS           | Oswaldo      | 85362        |
+-----+-----+-----+
3 rows in set (0.00 sec)
```

```
SELECT c.country_code, c.country_name, sum(p.area)
FROM country c JOIN province p ON c.country_code = p.country_code
GROUP BY c.country_code
ORDER BY sum(p.area) DESC;
```

5.

```
mysql> SELECT country_name
-> FROM country c JOIN province p JOIN city cit ON c.country_code = p.country_code AND p.country_code = cit.country_code
-> GROUP BY c.country_code
-> HAVING count(DISTINCT cit.city_name) >= @min_cities AND count(DISTINCT p.province_name) >= @min_provinces;
+-----+
| country_name |
+-----+
| Oswaldo      |
| Renlandia    |
+-----+
2 rows in set (0.00 sec)
```

```
set @min_cities = 5;
set @min_provinces = 1;
SELECT country_name
FROM country c JOIN province p JOIN city cit ON c.country_code =
p.country_code AND p.country_code = cit.country_code
GROUP BY c.country_code
HAVING count(DISTINCT cit.city_name) >= @min_cities AND count(DISTINCT
p.province_name) >= @min_provinces;
```

6.

```
mysql> set @gdp = 60000;
SELEQuery OK, 0 rows affected (0.00 sec)

mysql> SELECT c.country_code, c.gdp, sum(p.area)
-> FROM country c JOIN province p ON c.country_code = p.country_code
-> GROUP BY c.country_code
-> HAVING c.gdp >= @gdp
-> ORDER BY
-> CASE WHEN sum(p.area) <> 0 THEN sum(p.area) END DESC,
-> CASE WHEN c.gdp <> 0 THEN c.gdp END DESC;
+-----+-----+-----+
| country_code | gdp   | sum(p.area) |
+-----+-----+-----+
| GN           | 65000 | 126044       |
| OS           | 78000 | 85362        |
+-----+-----+-----+
2 rows in set (0.00 sec)
```

```
set @gdp = 60000;
SELECT c.country_code, c.gdp, sum(p.area)
FROM country c JOIN province p ON c.country_code = p.country_code
GROUP BY c.country_code
HAVING c.gdp >= @gdp
ORDER BY
CASE WHEN sum(p.area) <> 0 THEN sum(p.area) END DESC,
CASE WHEN c.gdp <> 0 THEN c.gdp END DESC;
```

7.

```
mysql> DROP VIEW IF EXISTS sym_borders;
Query OK, 0 rows affected (0.00 sec)

mysql> CREATE VIEW sym_borders AS
  -> SELECT *
  -> FROM border;
Query OK, 0 rows affected (0.01 sec)

mysql> INSERT INTO sym_borders
  -> SELECT country_code_2, country_code_1, border_length
  -> FROM border;
FROM sym_Query OK, 3 rows affected (0.00 sec)
Records: 3 Duplicates: 0 Warnings: 0

mysql>
mysql> SELECT * FROM sym_borders;
+-----+-----+-----+
| country_code_1 | country_code_2 | border_length |
+-----+-----+-----+
| GN             | OS             | 4800          |
| GN             | RL             | 4500          |
| OS             | GN             | 4800          |
| OS             | RL             | 1200          |
| RL             | GN             | 4500          |
| RL             | OS             | 1200          |
+-----+-----+-----+
6 rows in set (0.00 sec)
```

```
DROP VIEW IF EXISTS sym_borders;
CREATE VIEW sym_borders AS
SELECT *
FROM border;
INSERT INTO sym_borders
SELECT country_code_2, country_code_1, border_length
FROM border;

SELECT * FROM sym_borders;
```


8.

```
mysql> -- question 8
mysql> -- original
mysql> (SELECT c1.country_name, c2.country_name as c_gdp_high
-> FROM country c1 JOIN country c2 JOIN border bord
-> ON (bord.country_code_1 = c1.country_code AND bord.country_code_2 = c2.country_code)
-> WHERE c1.gdp < c2.gdp
-> AND c1.inflation > c2.inflation)
-> UNION
-> (SELECT c2.country_name, c1.country_name as c_gdp_high
-> FROM country c1 JOIN country c2 JOIN border bord
-> ON (bord.country_code_1 = c1.country_code AND bord.country_code_2 = c2.country_code)
-> WHERE c2.gdp < c1.gdp
-> AND c2.inflation > c1.inflation);
LECT +-----+
| country_name | c_gdp_high |
+-----+-----+
| Renlandia    | Geneva     |
| Renlandia    | Oswaldo    |
+-----+-----+
2 rows in set (0.00 sec)

mysql> -- new query
mysql> SELECT c1.country_name, c2.country_name as c_gdp_high
-> FROM country c1 JOIN country c2 JOIN sym_borders bord
-> ON (bord.country_code_1 = c1.country_code AND bord.country_code_2 = c2.country_code)
-> WHERE (c1.gdp < c2.gdp
-> AND c1.inflation > c2.inflation)
-> OR (c2.gdp < c1.gdp
-> AND c2.inflation > c1.inflation);
+-----+-----+
| country_name | c_gdp_high |
+-----+-----+
| Renlandia    | Geneva     |
| Renlandia    | Oswaldo    |
| Geneva       | Renlandia  |
| Oswaldo      | Renlandia  |
+-----+-----+
4 rows in set (0.00 sec)
```

-- question 8

-- original

```
(SELECT c1.country_name, c2.country_name as c_gdp_high
FROM country c1 JOIN country c2 JOIN border bord
ON (bord.country_code_1 = c1.country_code AND bord.country_code_2 =
c2.country_code)
```

WHERE c1.gdp < c2.gdp

AND c1.inflation > c2.inflation)

UNION

```
(SELECT c2.country_name, c1.country_name as c_gdp_high
```

FROM country c1 JOIN country c2 JOIN border bord

```
ON (bord.country_code_1 = c1.country_code AND bord.country_code_2 =
c2.country_code)
```

WHERE c2.gdp < c1.gdp

AND c2.inflation > c1.inflation);

-- new query

```
SELECT c1.country_name, c2.country_name as c_gdp_high
```

```
FROM country c1 JOIN country c2 JOIN sym_borders bord
```

```

ON (bord.country_code_1 = c1.country_code AND bord.country_code_2 =
c2.country_code)
WHERE (c1.gdp < c2.gdp
AND c1.inflation > c2.inflation)
OR (c2.gdp < c1.gdp
AND c2.inflation > c1.inflation);

```

9.

```

mysql> -- question 9
mysql> SELECT c1.country_name, avg(c2.gdp), avg(c2.inflation)
-> FROM country c1 JOIN country c2 JOIN sym_borders b
-> ON b.country_code_1 = c1.country_code
-> AND b.country_code_2 = c2.country_code
-> GROUP BY b.country_code_1
-> ORDER BY
-> CASE WHEN avg(c2.gdp) <> 0 THEN avg(c2.gdp) END ASC,
-> CASE WHEN avg(c2.inflation) <> 0 THEN avg(c2.inflation) END ASC;
+-----+-----+-----+
| country_name | avg(c2.gdp) | avg(c2.inflation) |
+-----+-----+-----+
| Oswaldo      | 59500.0000   | 4.8999998569488525 |
| Geneva       | 66000.0000   | 7.199999809265137  |
| Renlandia    | 71500.0000   | 4.3999998569488525 |
+-----+-----+-----+
3 rows in set (0.01 sec)

```

```

-- question 9
SELECT c1.country_name, avg(c2.gdp), avg(c2.inflation)
FROM country c1 JOIN country c2 JOIN sym_borders b
ON b.country_code_1 = c1.country_code
AND b.country_code_2 = c2.country_code
GROUP BY b.country_code_1
ORDER BY
CASE WHEN avg(c2.gdp) <> 0 THEN avg(c2.gdp) END ASC,
CASE WHEN avg(c2.inflation) <> 0 THEN avg(c2.inflation) END ASC;

```


10.

```
mysql> -- question 10
mysql> -- part 1: Show all cities which are in a country with a designated relationship between
mysql> -- the country's inflation and gdp, and those of a bordering country. Should return
mysql> -- the name of the city, name of the province the city is in, and the population of the
mysql> -- city, ordered by population of the city going from highest to lowest.
mysql> -- first argument: country 1 has a higher gdp, higher inflation
mysql> SELECT c.city_name, c.province_name, c.population
-> FROM country c1 JOIN country c2 JOIN sym_borders b JOIN city c
-> ON b.country_code_1 = c1.country_code
-> AND b.country_code_2 = c2.country_code
-> AND c.country_code = c1.country_code
-> WHERE c1.gdp > c2.gdp
-> AND c1.inflation > c2.inflation
-> ORDER BY c.population DESC;
+-----+-----+-----+
| city_name | province_name | population |
+-----+-----+-----+
| Slaren    | Oslo         | 590190    |
| Belogonia | St. Janice   | 97635     |
| Britano   | St. Janice   | 33434     |
| Sluurgan  | Oslo         | 5919      |
| Tesa      | Oslo         | 964       |
+-----+-----+-----+
5 rows in set (0.01 sec)
```

```
mysql> -- second argument: country 1 has a higher gdp, lower inflation
mysql> SELECT c.city_name, c.province_name, c.population
-> FROM country c1 JOIN country c2 JOIN sym_borders b JOIN city c
-> ON b.country_code_1 = c1.country_code
-> AND b.country_code_2 = c2.country_code
-> AND c.country_code = c1.country_code
-> WHERE c1.gdp < c2.gdp
-> AND c1.inflation > c2.inflation
-> ORDER BY c.population DESC;
```

```
+-----+-----+-----+
| city_name | province_name | population |
+-----+-----+-----+
| Gonba     | Flaubury     | 53146     |
| Gonba     | Flaubury     | 53146     |
| Whita     | Flaubury     | 52743     |
| Whita     | Flaubury     | 52743     |
| Stombus   | Huport       | 49384     |
| Stombus   | Huport       | 49384     |
| Piolas    | Huport       | 46626     |
| Piolas    | Huport       | 46626     |
| Outling   | Flaubury     | 127       |
| Outling   | Flaubury     | 127       |
+-----+-----+-----+
10 rows in set (0.00 sec)
```

```
mysql>
mysql> -- part 2: Show all provinces with a designated relationship between the area of that province
mysql> -- and that of a province in a bordering country. Should return the name of each province,
mysql> -- the area of each province, and the names of the countries that border each other.
mysql> -- first argument: province 1 has a higher area
mysql> SELECT p1.province_name, p2.province_name, p1.area, p2.area, c1.country_code, c2.country_code
-> FROM country c1 JOIN country c2 JOIN sym_borders b
-> JOIN province p1 JOIN province p2
-> ON b.country_code_1 = c1.country_code
-> AND b.country_code_2 = c2.country_code
-> AND c1.country_code = p1.country_code
-> AND c2.country_code = p2.country_code
-> WHERE p1.area > p2.area
-> ORDER BY p1.area DESC;
```

```
sec+-----+-----+-----+-----+-----+-----+
| province_name | province_name | area | area | country_code | country_code |
+-----+-----+-----+-----+-----+-----+
| Huport       | Antalens     | 91030 | 72003 | RL           | GN           |
| Huport       | Huport       | 91030 | 54041 | RL           | GN           |
| Huport       | Oslo         | 91030 | 8712  | RL           | OS           |
| Huport       | St. Janice   | 91030 | 76650 | RL           | OS           |
| St. Janice   | Antalens     | 76650 | 72003 | OS           | GN           |
| St. Janice   | Flaubury     | 76650 | 5690  | OS           | RL           |
| St. Janice   | Huport       | 76650 | 54041 | OS           | GN           |
| Antalens     | Flaubury     | 72003 | 5690  | GN           | RL           |
| Antalens     | Oslo         | 72003 | 8712  | GN           | OS           |
| Huport       | Flaubury     | 54041 | 5690  | GN           | RL           |
| Huport       | Oslo         | 54041 | 8712  | GN           | OS           |
| Oslo         | Flaubury     | 8712  | 5690  | OS           | RL           |
+-----+-----+-----+-----+-----+-----+

```

```
mysql> -- second argument: province 1 has a lower area
mysql> SELECT p1.province_name, p2.province_name, p1.area, p2.area, c1.country_code, c2.country_code
-> FROM country c1 JOIN country c2 JOIN sym_borders b
-> JOIN province p1 JOIN province p2
-> ON b.country_code_1 = c1.country_code
-> AND b.country_code_2 = c2.country_code
-> AND c1.country_code = p1.country_code
-> AND c2.country_code = p2.country_code
-> WHERE p1.area < p2.area
-> ORDER BY p1.area DESC;
```

province_name	province_name	area	area	country_code	country_code
St. Janice	Huport	76650	91030	OS	RL
Antalens	Huport	72003	91030	GN	RL
Antalens	St. Janice	72003	76650	GN	OS
Huport	Huport	54041	91030	GN	RL
Huport	St. Janice	54041	76650	GN	OS
Oslo	Huport	8712	91030	OS	RL
Oslo	Antalens	8712	72003	OS	GN
Oslo	Huport	8712	54041	OS	GN
Flaubury	Antalens	5690	72003	RL	GN
Flaubury	Huport	5690	54041	RL	GN
Flaubury	Oslo	5690	8712	RL	OS
Flaubury	St. Janice	5690	76650	RL	OS

```
12 rows in set (0.00 sec)

mysql>
```

-- question 10

-- part 1: Show all cities which are in a country with a designated relationship between

-- the country's inflation and gdp, and those of a bordering country. Should return

-- the name of the city, name of the province the city is in, and the population of the

-- city, ordered by population of the city going from highest to lowest.

-- first argument: country 1 has a higher gdp, higher inflation

```
SELECT c.city_name, c.province_name, c.population
```

```
FROM country c1 JOIN country c2 JOIN sym_borders b JOIN city c
```

```
ON b.country_code_1 = c1.country_code
```

```
AND b.country_code_2 = c2.country_code
```

```
AND c.country_code = c1.country_code
```

```
WHERE c1.gdp > c2.gdp
```

```
AND c1.inflation > c2.inflation
```

```
ORDER BY c.population DESC;
```

-- second argument: country 1 has a higher gdp, lower inflation

```
SELECT c.city_name, c.province_name, c.population
```

```
FROM country c1 JOIN country c2 JOIN sym_borders b JOIN city c
```

```
ON b.country_code_1 = c1.country_code
```

```
AND b.country_code_2 = c2.country_code
```

```
AND c.country_code = c1.country_code
```

```
WHERE c1.gdp < c2.gdp
```

```
AND c1.inflation > c2.inflation
```

```
ORDER BY c.population DESC;
```

```
-- part 2: Show all provinces with a designated relationship between  
the area of that province
```

```
--           and that of a province in a bordering country. Should  
return the name of each province,
```

```
--           the area of each province, and the names of the countries  
that border each other.
```

```
-- first argument: province 1 has a higher area
```

```
SELECT p1.province_name, p2.province_name, p1.area, p2.area,  
c1.country_code, c2.country_code
```

```
FROM country c1 JOIN country c2 JOIN sym_borders b
```

```
JOIN province p1 JOIN province p2
```

```
ON b.country_code_1 = c1.country_code
```

```
AND b.country_code_2 = c2.country_code
```

```
AND c1.country_code = p1.country_code
```

```
AND c2.country_code = p2.country_code
```

```
WHERE p1.area > p2.area
```

```
ORDER BY p1.area DESC;
```

```
-- second argument: province 1 has a lower area
```

```
SELECT p1.province_name, p2.province_name, p1.area, p2.area,  
c1.country_code, c2.country_code
```

```
FROM country c1 JOIN country c2 JOIN sym_borders b
```

```
JOIN province p1 JOIN province p2
```

```
ON b.country_code_1 = c1.country_code
```

```
AND b.country_code_2 = c2.country_code
```

```
AND c1.country_code = p1.country_code
```

```
AND c2.country_code = p2.country_code
```

```
WHERE p1.area < p2.area
```

```
ORDER BY p1.area DESC;
```